

Lab C

CPU used:

- AMD Ryzen 5 5500U
 - 6 core/12 threads

For all the benchmarks we used: futhark bench, to do our benchmarks

Exercise 1

Testing process and process_idx

Given two input lists $s1$ & $s2$ such that:

$s1 = [23, 45, -23, 44, 23, 54, 23, 12, 34, 54, 7, 2, 4, 67]$

$s2 = [-2, 3, 4, 57, 34, 2, 5, 56, 56, 3, 3, 5, 77, 89]$

process outputs:

73i32

and process_idx outputs:

73i32

12i64

Benchmarking results

By using the multicore backend we can see that the speedup (compared to when using the c backend) scales with the input size for both `process` and `process_idx`. As can be seen in *Fig 1* and *Fig 2*, both `process` and `process_idx` scale in a similar manner, that is, both scale linearly.

Graphs

Data generated by futhark:

two_%_i32s:

futhark dataset -b --i32-bounds=-10000:10000 -g [*]i32 -g [*]i32 > \$@

With the sizes :

100, 1000, 10000, 100000, 1000000, 5000000, 10000000




Fig 1: process runtime and speed up on multicore compared to c




Fig 2: process_idx runtime and speedup on multicore compared to c

Exercise 2

Proof of left-identity

Here we prove that $(0, \text{false})$ is a left-neutral element of:

$$(v1, f1) \oplus (v2, f2) = (\text{if } f2 \text{ then } v2 \text{ else } v1 \oplus v2, f1 \vee f2)$$

$$\begin{aligned}
 (0, \text{false}) \oplus (v, f) &= (\text{if } f \text{ then } v \text{ else } 0 \oplus v, \text{false} \vee f) \quad (0 \text{ is neutral for } \oplus \text{ and } \text{false} \text{ is neutral for } \vee) \\
 &= (\text{if } f \text{ then } v \text{ else } v, f) \\
 &= (v, f)
 \end{aligned}$$

■

Benchmarking results for segscan and segreduce

The function we benchmark is addition (+).

When benchmarking `segscan`, `segreduce`, `scan` and `reduce` we used both the `c` backend and the multicore backend. In *Fig 4* and *Fig 5* we can see how our `segscan` and how the built-in `scan` compares and in *Fig 8* and *Fig 9* we can see how our `segreduce` compares to the built-in `reduce`. In both cases we get slowdowns which is not surprising since more work is done in `segscan` and `segreduce`.

Both `segscan` and `scan` have the same scaling. This is also the case for `segreduce` and `reduce`: That we have linear scaling but `reduce` is just an order of magnitude faster, which makes it look like constant scaling in *Fig 8* and *9*. One reason for this may be that we use addition as the operator, because of this `futhark` uses `reduce_comm`¹ instead of `reduce` when benchmarking and `reduce_comm` is faster. However, both scale linearly with the input, which can be seen in *Fig 6* and *Fig 7* (note log scaled x-axis).

¹ <https://futhark-lang.org/docs/prelude/doc/prelude/soacs.html#4953>

Graphs

Data generated by futhark:

For built-in:

one_%_i32s:

```
futhark dataset -b --i32-bounds=-10000:10000 -g [$*]i32 > $@
```

For segmented:

i32_%_bools:

```
futhark dataset -b --i32-bounds=-10000:10000 -g [$*]i32 -g [$*]bool > $@
```

Both with the sizes :

100, 1000, 10000, 100000, 1000000, 5000000, 10000000

Segscan




Fig 3: segscan runtime and speed up on multicore compared to c




Fig 4: segscan compared to scan on c

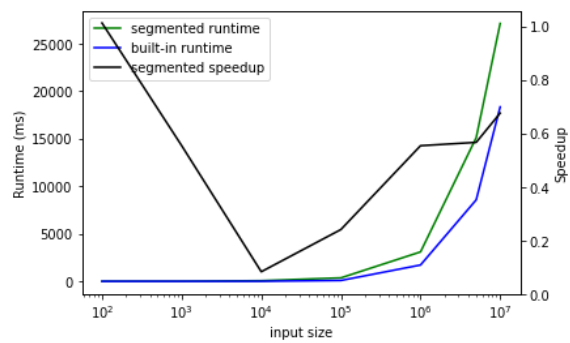


Fig 5: segscan compared to scan on multicore

Segreduce




Fig 6: segreduce runtime and speedup on multicore compared to c

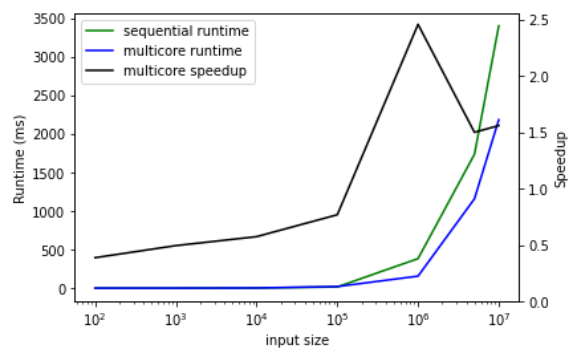


Fig 7: reduce runtime and speedup on multicore compared to c

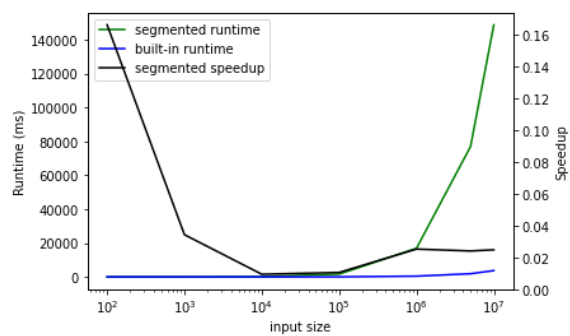


Fig 8: segreduce compared to reduce on c

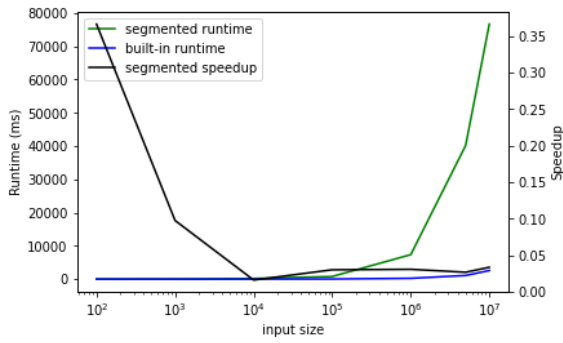


Fig 9: *segreduce* compared to *reduce* on multicore

Asymptotic complexity of our `reduce_by_index`

The two most expensive operations it uses are `radix_sort_by_key` and `segreduce`. The work and span complexity for `radix_sort_by_key` is $O(k n)$ and $O(k \log n)^2$ respectively, where k is constant. To get the complexity of `segreduce` we first need the complexity of `segscan`.

Both `segscan` and `scan` have the same complexity. The complexity of `scan` (which is the only operation used in `segscan`) has a work and span complexity that depends on the operator used in the function. In the documentation the work is written as $O(n W(op))$ and the span is written as $O(\log(n) W(op))$ where $W(op)$ is the work complexity for the operator³.

The most expensive operation of `segreduce` is `segscan` and thus the complexity for `segreduce` is the same as the complexity for `segscan`.

Going back to the work and span complexity of `radix_sort_by_key`; Since k is constant and $W(op)$ can be of any complexity we can conclude that the complexity for our `reduce_by_index` is:

- Work: $O(n W(op))$
- Span: $O(\log(n) W(op))$

If we compare it to the built-in `reduce_by_index` who has⁴:

- Work: $O(n W(op))$
- Span:
 - Worst case: $O(n W(op))$
 - Best case: $O(W(op))$

Benchmarking results for `reduce_by_index`

The function we benchmark is addition (+).

²

https://futhark-lang.org/pkgs/github.com/diku-dk/sorts/0.4.1/doc/lib/github.com/diku-dk/sorts/radix_sort.html#abstract

³ <https://futhark-lang.org/docs/prelude/doc/prelude/soacs.html#5011>

⁴ <https://futhark-lang.org/docs/prelude/doc/prelude/soacs.html#4967>

Our `reduce_by_index` worked by first sorting and then doing a segmented reduction. This should in practice be better to do than the built-in `reduce_by_index` when we have many millions of bins⁵. But in practice, our solution is much slower than the built-in variant.

As with `segreduce`, our function scales like the built-in one, it is just a lot slower. We can see, when comparing *Fig 10* and *Fig 11*, that both our implementation and the built-in one scales linearly.

Data generated by futhark:

one_%_two:

```
futhark dataset -b --i32-bounds=-10000:10000 \
                --i64-bounds=-10:$* \
                -g [$*]i32 -g [$*]i64 -g [$*]i32 > $@
```

With the sizes :

100, 1000, 10000, 100000, 1000000, 5000000, 10000000

Graphs

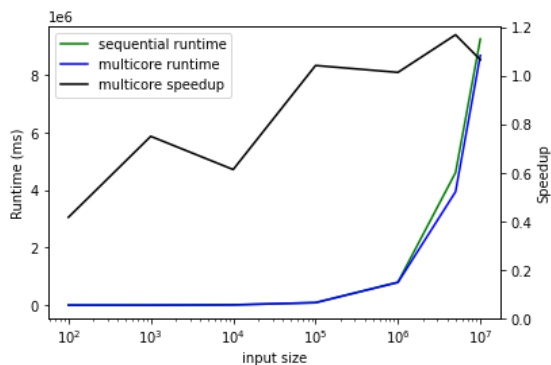


Fig 10: reduce_by_index runtime and speedup on multicore compared to c




Fig 11: built-in reduce_by_index runtime and speedup on multicore compared to c

⁵ <https://futhark.readthedocs.io/en/latest/performance.html#histograms>




Fig 12: `reduce_by_index` compared to built-in `reduce_by_index` on *c*

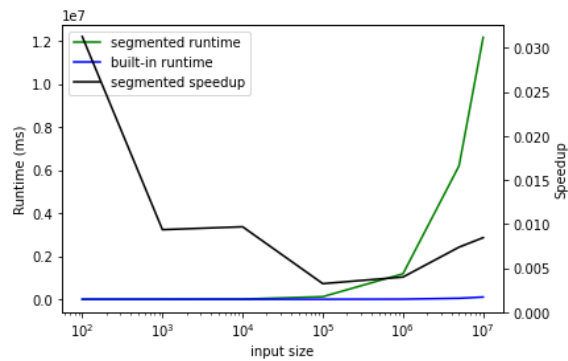


Fig 13: `reduce_by_index` compared to built-in `reduce_by_index` on multicore

Exercise 3

Benchmarking results

<i>parameter</i>	<i>default value</i>
<i>h,w</i>	400
<i>n</i>	20
<i>abs_temp</i>	0.5
<i>samplerate</i>	0.1

<i>parameters</i>	<i>sizes we test</i>
<i>w,h,h and w</i>	100, 200, 400, 800, 1600
<i>n</i>	10,20,40,80

Graphs




Fig 14: x-axis: size of h




Fig 15: x-axis: size of w

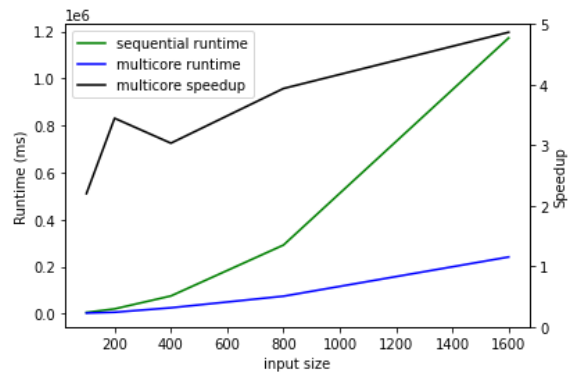


Fig 16: x-axis: size of w and h




Fig 17: x-axis: size of n

Conclusion

We got linear scaling when we varied either w , h or n and quadratic when we varied both w and h at the same time, as one would expect. We can also see that we got a good speedup (between 3-5x) and this is without thinking about creating a parallel program. We just needed to change the backend to multicore and our solution was parallized.




Fig 18: visualization of the ising model